

# Basics of Reinforcement Learning

# Overview



**HMMs** The probability transformation matrix  $a_{ij}$  allows forward and backward propagation

**Dynamic programming** An objective function and dynamics uses backward propagation

**Reinforcement Learning** A reward function and probabilistic dynamics uses forward propagation with iteration

**Model-free Reinforcement Learning** A reward function and probabilistic dynamics uses forward propagation with search for model estimation

# The Urn model of a HMM



# Dynamic Programming

Final value problem



The dynamics is expressed by a difference equation,

$$\mathbf{x}(k+1) = f[\mathbf{x}(k), \mathbf{u}(k)]$$

The initial condition is:

$$\mathbf{x}(0) = \mathbf{x}_0$$

The allowable control is also discrete:

$$\mathbf{u}(k) \in U, k = 0, \dots, T$$

The integral in the objective function is expressed as a sum:

$$J = \psi[\mathbf{x}(T)] + \sum_{k=0}^{T-1} \ell[\mathbf{u}(k), \mathbf{x}(k)]$$

# Solution: discretize the state space and work backwards

Let  $V(k)$  keep track  
of estimates  
of the  
objective function  $J$

At the end. when  $k=T$   $V(\mathbf{x}, N) = \psi[\mathbf{x}(T)]$


One step back,

$$V(\mathbf{x}, T-1) = \max_{\mathbf{u} \in U} \{ \ell[\mathbf{u}(T-1), \mathbf{x}(T-1)] + \psi[\mathbf{x}(T)] \}$$


And in general, for  $k < T-1$ ,

$$V(\mathbf{x}, k-1) = \max_{\mathbf{u} \in U} \{ \ell[\mathbf{u}(k-1), \mathbf{x}(k-1)] + V[\mathbf{x}(k), k] \}, k = 0, \dots, T$$

Discretizing the state  
to a fine scale is very expensive.  
*Bellman*, the originator of  
Dynamic Programming  
called it  
**“The curse of dimensionality”**



# Basic RL Problems



Location of reward  
uncertain

Transitions between  
states  
uncertain

Policy constantly  
changing



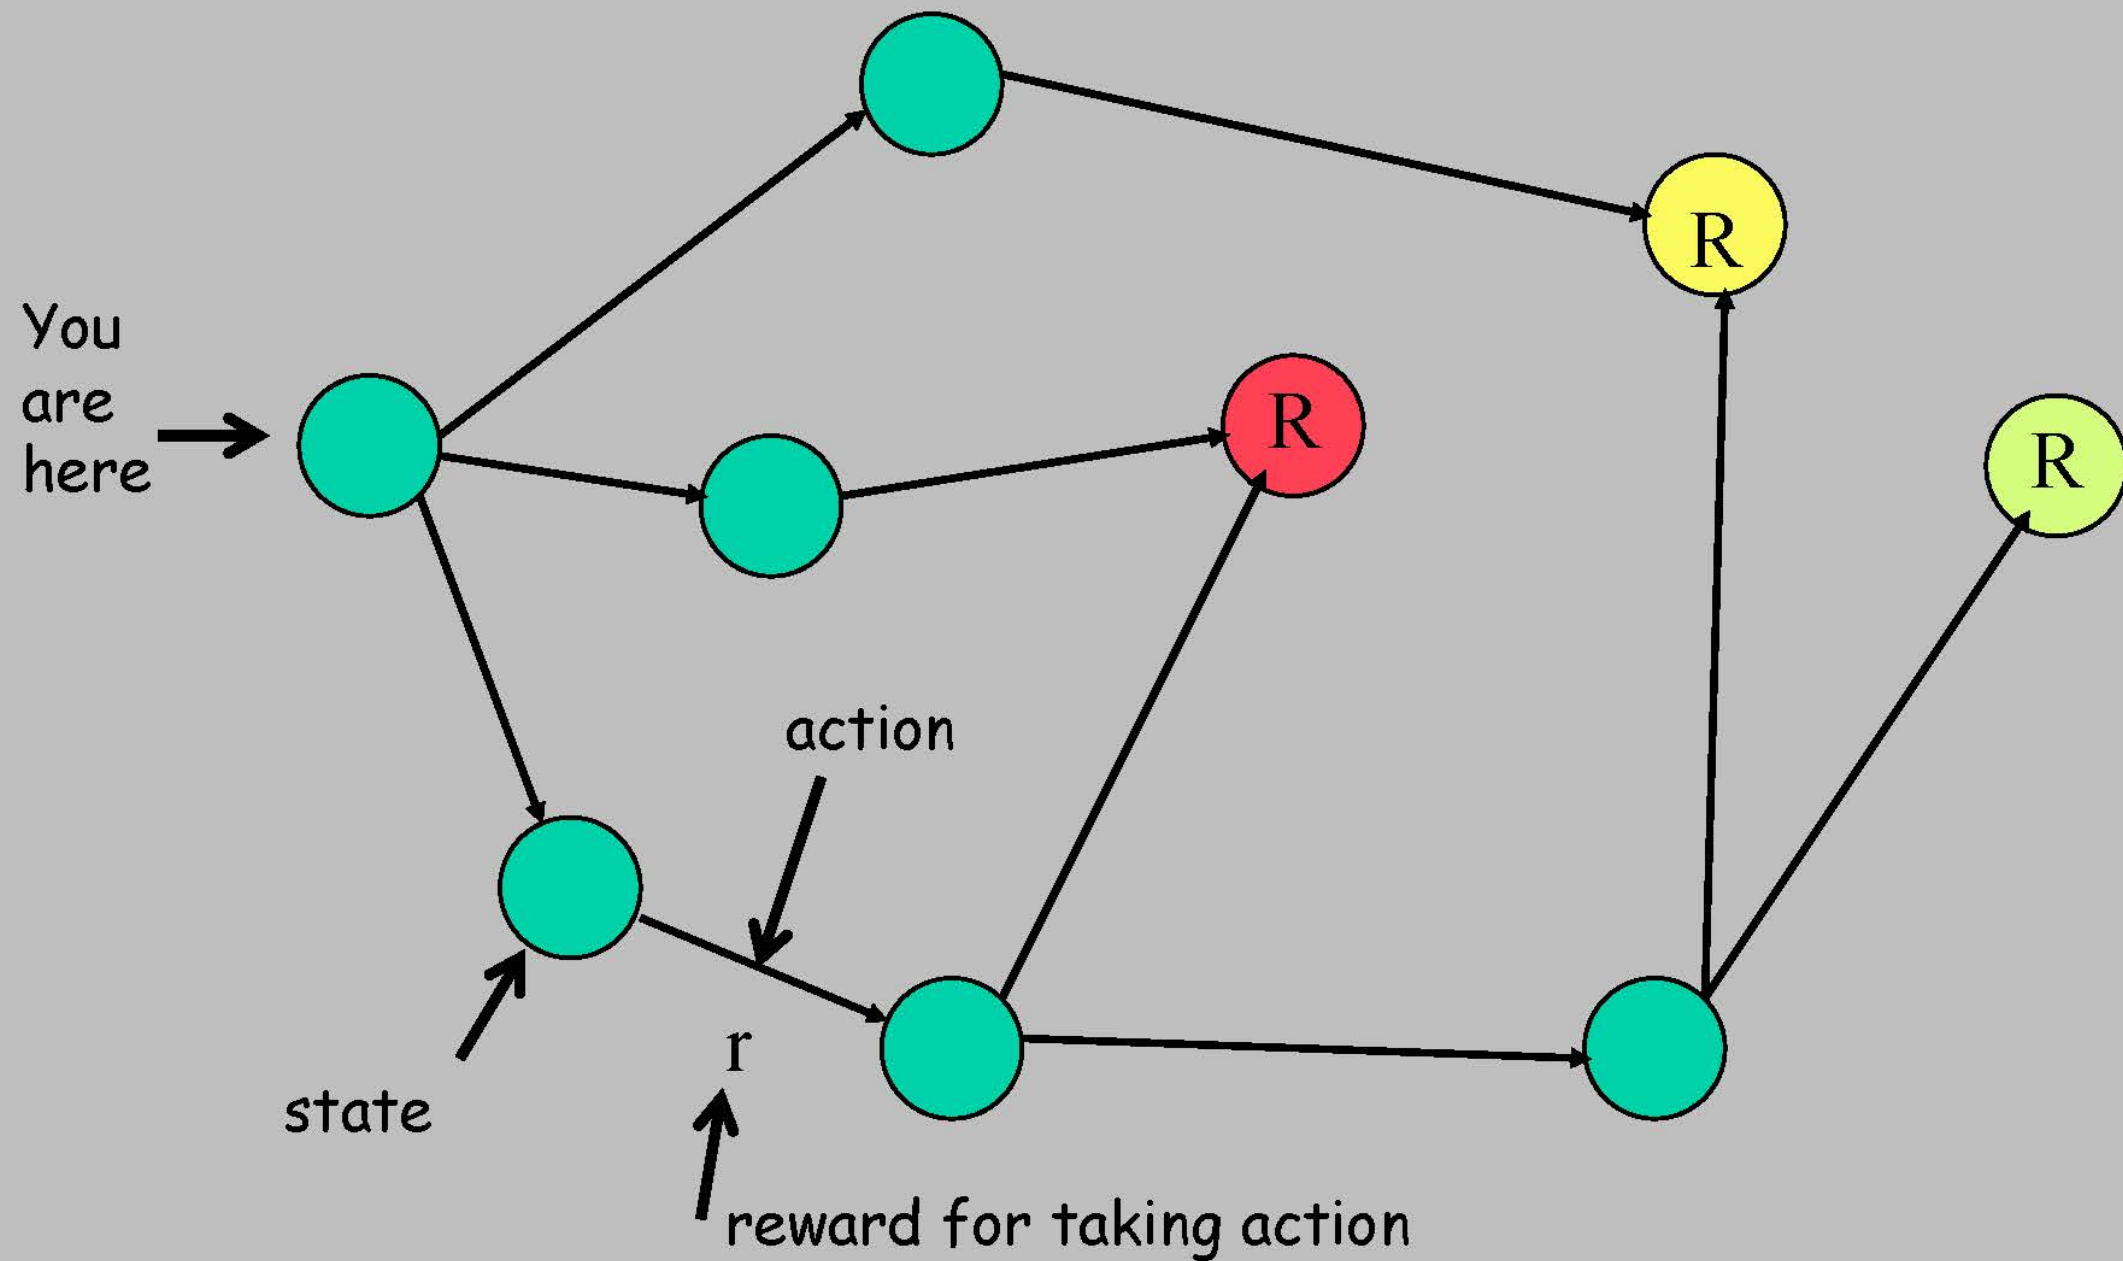
# Markov Decision Processes

Problems with delayed reinforcement are well modeled as *Markov decision processes* (MDPs). An MDP consists of

- a set of states  $\mathcal{S}$  ,
- a set of actions  $\mathcal{A}$  ,
- a reward function  $R : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$  , and
- a state transition function  $T : \mathcal{S} \times \mathcal{A} \rightarrow \Pi(\mathcal{S})$  , where a member of  $\Pi(\mathcal{S})$  is a probability distribution over the set  $\mathcal{S}$  (i.e. it maps states to probabilities). We write  $T(s,a,s')$  for the probability of making a transition from state  $s$  to state  $s'$  using action  $a$ .

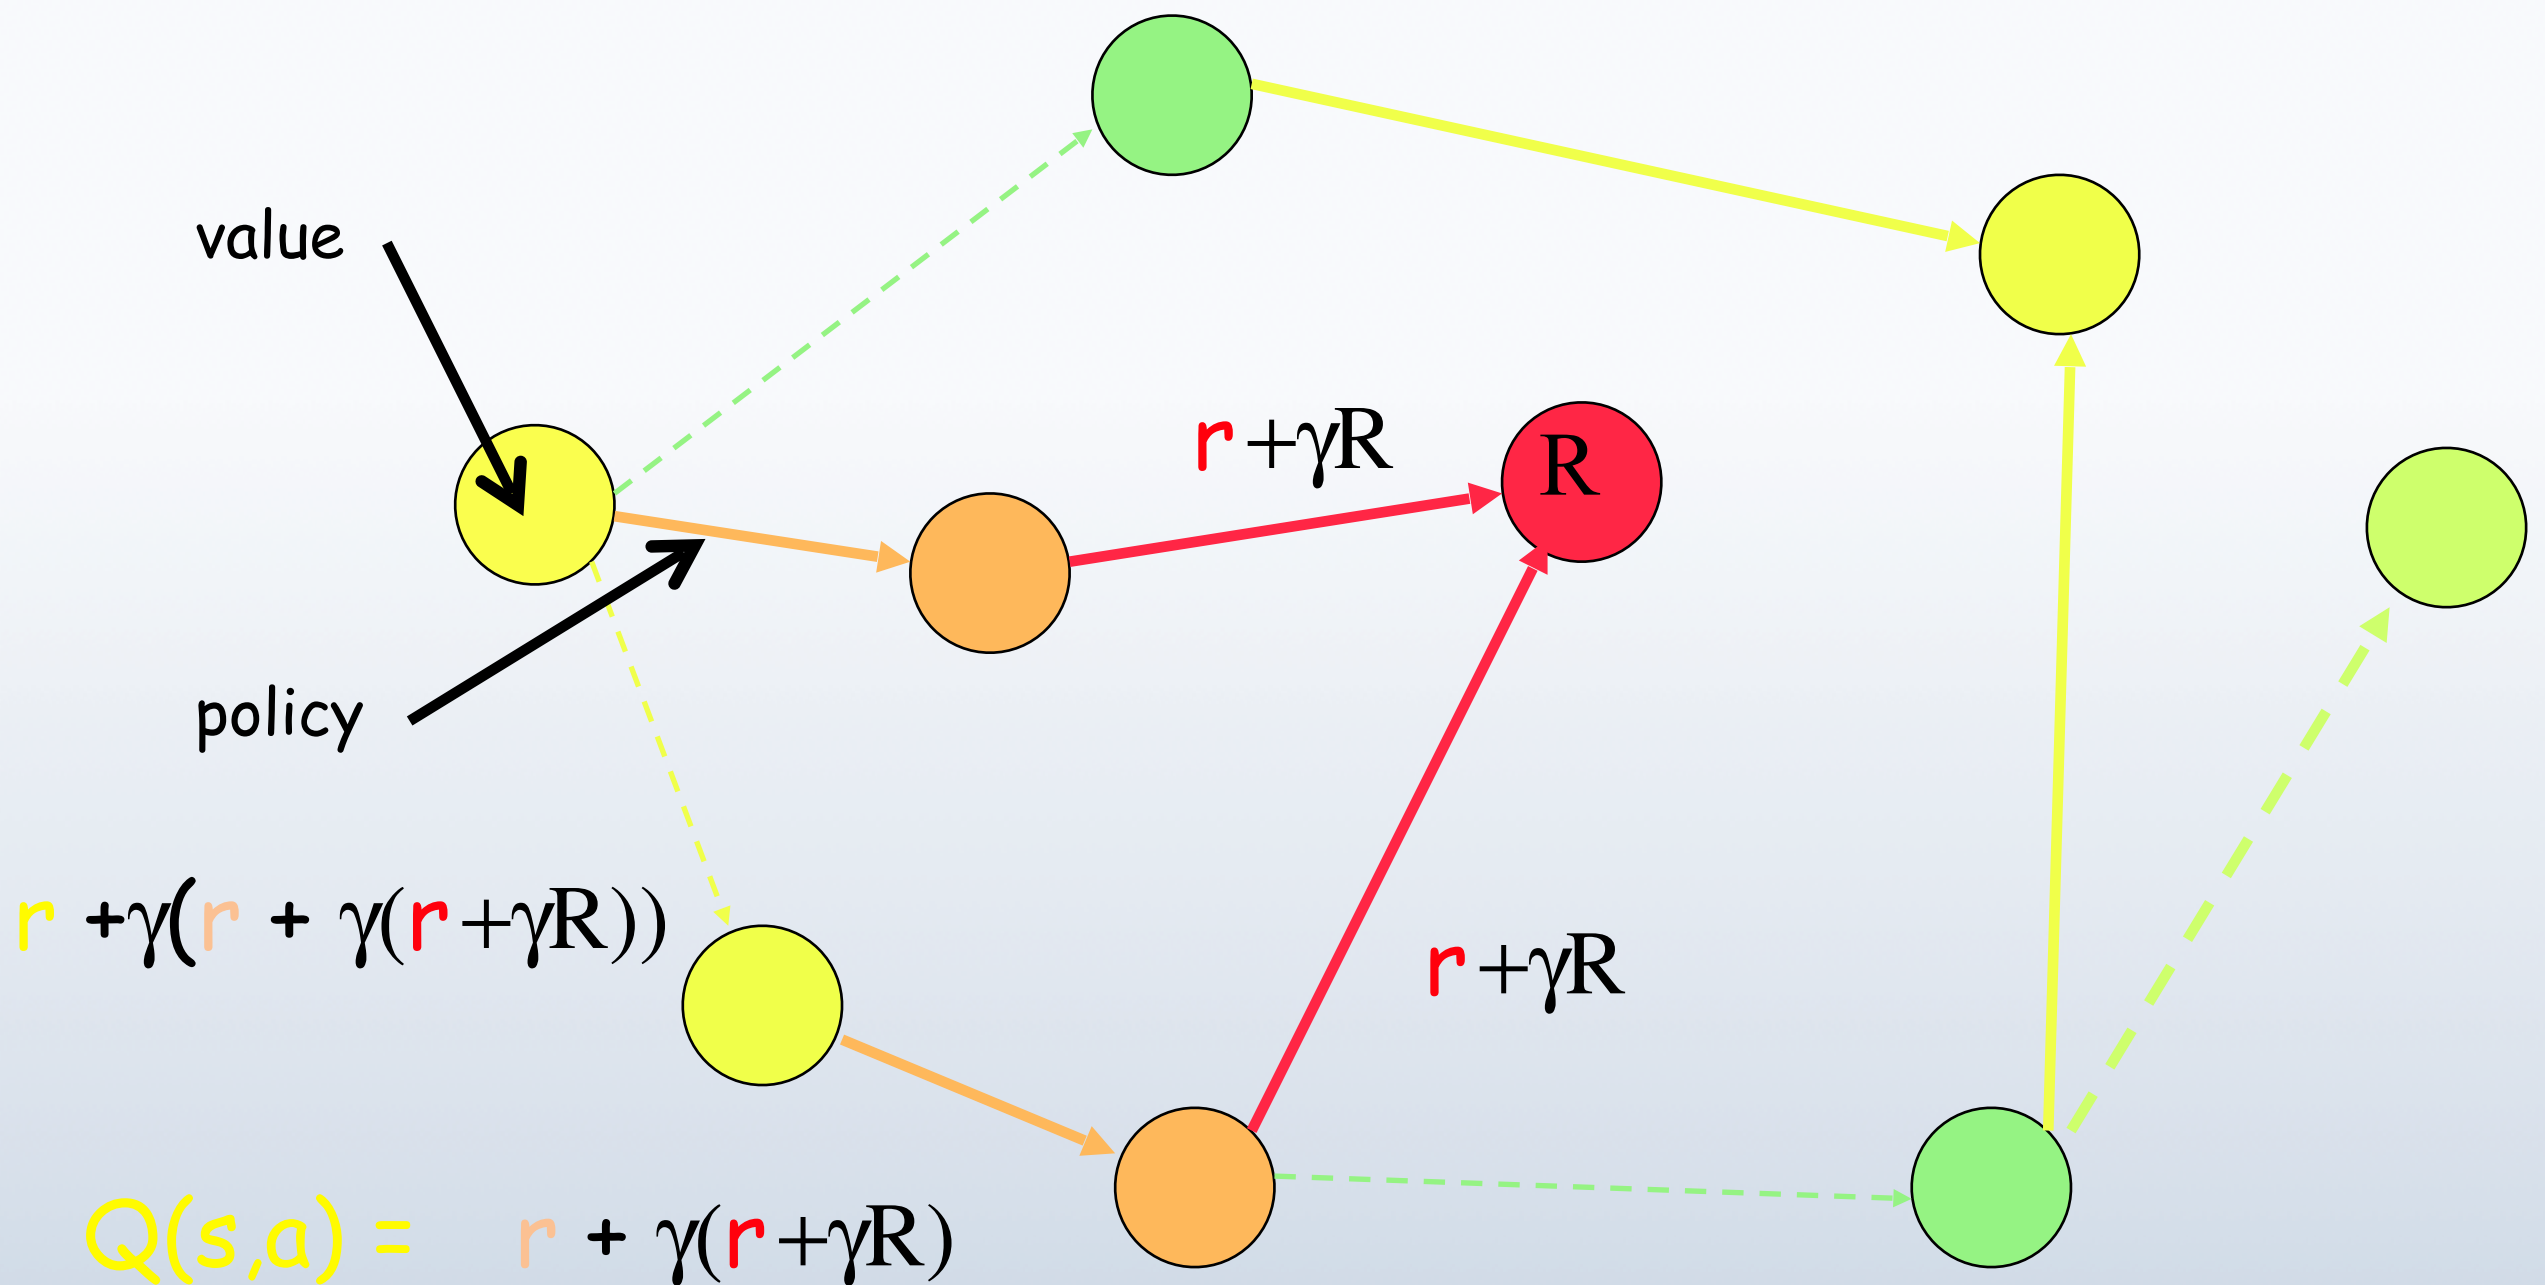
The state transition function probabilistically specifies the next state of the environment as a function of its current state and the agent's action. The reward function specifies expected instantaneous reward as a function of the current state and action. The model is *Markov* if the state transitions are independent of any previous environment states or agent actions. There are many good references to MDP models

# Reinforcement Learning Primer : Before Learning



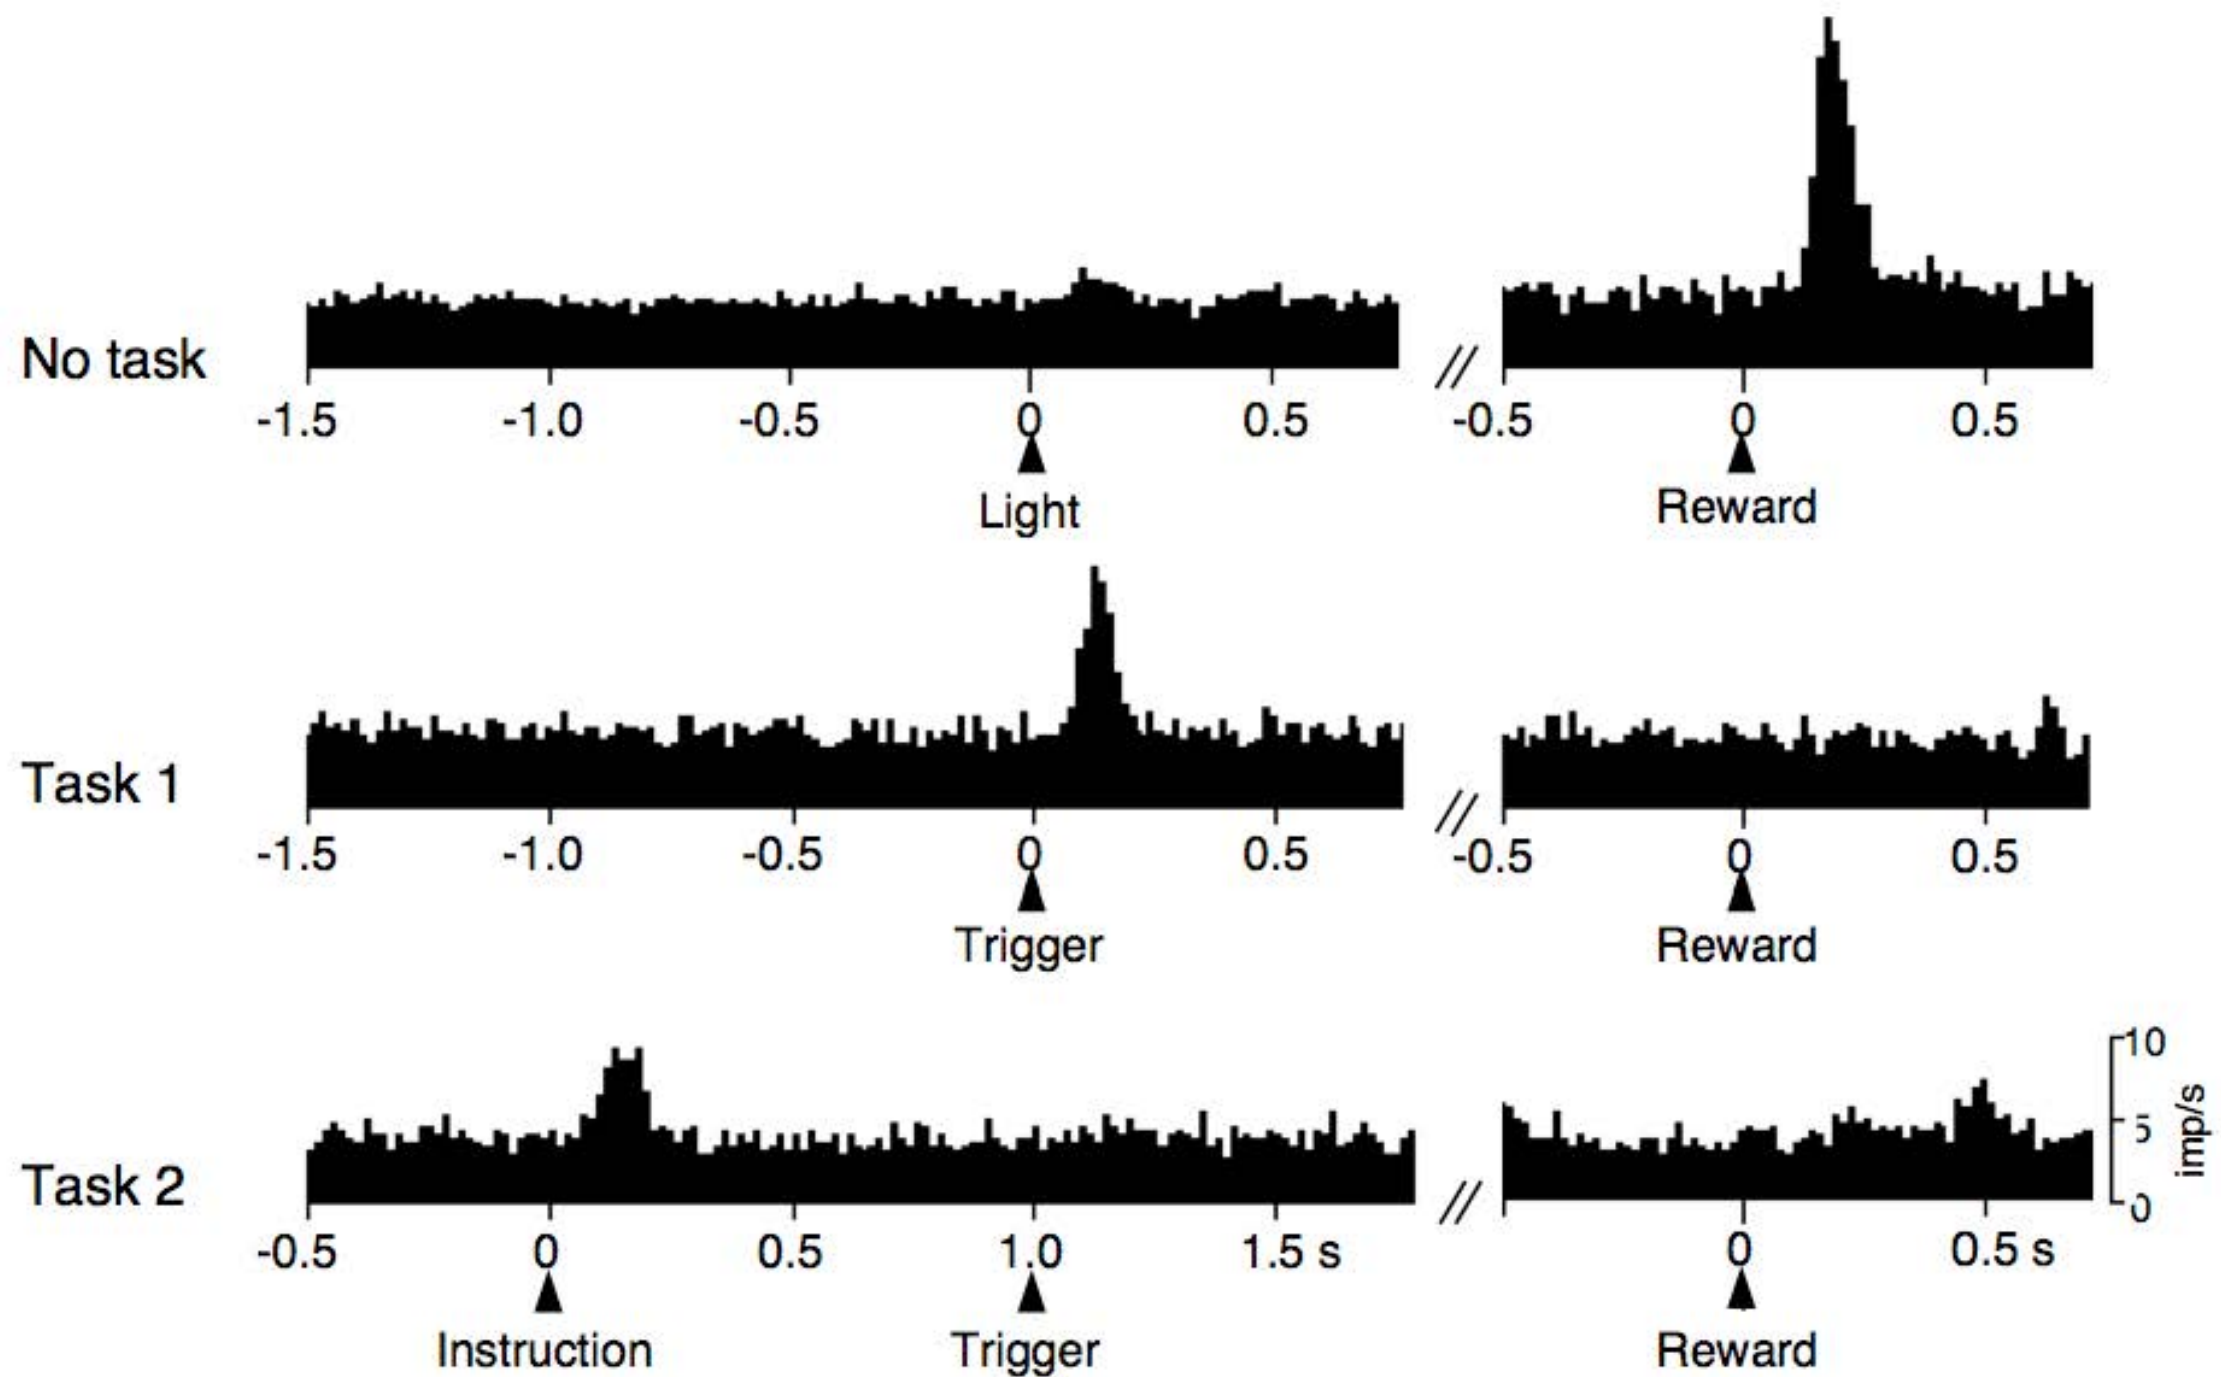


# Reinforcement Learning Primer




By trying different actions from different starting points, gradually learn the expected reward value from any starting point

# A Monkey uses Secondary Reward




Do dopamine neurons report an error  
in the prediction of reward?


No prediction  
Reward occurs



Reward predicted  
Reward occurs



Reward predicted  
No reward occurs



# The basic Reinforcement Learning model

We will speak of the optimal *value* of a state--it is the expected infinite discounted sum of reward that the agent will gain if it starts in that state and executes the optimal policy. Using  $\pi$  as a complete decision policy, it is written

$$V^*(s) = \max_{\pi} E \left( \sum_{t=0}^{\infty} \gamma^t r_t \right) .$$

This optimal value function is unique and can be defined as the solution to the simultaneous equations

$$V^*(s) = \max_a \left( R(s, a) + \gamma \sum_{s' \in S} T(s, a, s') V^*(s') \right) , \forall s \in S , \quad (1)$$

which assert that the value of a state  $s$  is the expected instantaneous reward plus the expected discounted value of the next state, using the best available action. Given the optimal value function, we can specify the optimal policy as

$$\pi^*(s) = \arg \max_a \left( R(s, a) + \gamma \sum_{s' \in S} T(s, a, s') V^*(s') \right) .$$

# Value Iteration

```
initialize  $V(s)$  arbitrarily
```

```
loop until policy good enough
```

```
  loop for  $s \in \mathcal{S}$ 
```

```
    loop for  $a \in \mathcal{A}$ 
```

$$Q(s, a) := R(s, a) + \gamma \sum_{s' \in \mathcal{S}} T(s, a, s') V(s')$$

$$V(s) := \max_a Q(s, a)$$

```
  end loop
```

```
end loop
```



# Policy Iteration

choose an arbitrary policy  $\pi'$

loop

$\pi := \pi'$

compute the value function of policy  $\pi$  :

solve the linear equations

$$V_{\pi}(s) = R(s, \pi(s)) + \gamma \sum_{s' \in S} T(s, \pi(s), s') V_{\pi}(s')$$

improve the policy at each state:

$$\pi'(s) := \arg \max_a \left( R(s, a) + \gamma \sum_{s' \in S} T(s, a, s') V_{\pi}(s') \right)$$

until  $\pi = \pi'$



# Temporal Difference Learning

$$\bar{V}_t = \sum_{i=0}^{\infty} \gamma^i r_{t+i}$$

where  $0 \leq \gamma < 1$ . This formula can be expanded

$$\bar{V}_t = r_t + \sum_{i=1}^{\infty} \gamma^i r_{t+i}$$

by changing the index of  $i$  to start from 0.

$$\bar{V}_t = r_t + \sum_{i=0}^{\infty} \gamma^{i+1} r_{t+i+1}$$

$$\bar{V}_t = r_t + \gamma \sum_{i=0}^{\infty} \gamma^i r_{t+i+1}$$

$$\bar{V}_t = r_t + \gamma \bar{V}_{t+1}$$

Thus, the reinforcement is the difference between the ideal prediction and the current prediction.

$$r_t = \bar{V}_t - \gamma \bar{V}_{t+1}$$

# Q - Learning

Temporal difference learning [Sutton and Barto, 1998], uses the error between the current estimated values of states and the observed reward to drive learning. In a related Q-learning form, the estimate of the quality value of a state-action pair is adjusted by this error  $\delta_Q$  using a learning rate  $\alpha$ :

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \delta_Q \quad (3)$$

Two fundamental learning rules for  $\delta_Q$  are 1) the original Q-learning rule [Watkins, 1989] and 2) SARSA [Rummery and Niranjan, 1994]. While Q-learning rule is an off-policy rule, i.e. it uses errors between current observations and estimates of the values for following an optimal policy, while actually following a potentially suboptimal policy during learning, SARSA<sup>1</sup> is an on-policy learning rule, i.e. the updates of the state and action values reflect the current policy derived from these value estimates. While in the general case of Q-learning, the temporal difference is:

$$\delta_Q = r_t + \gamma \max_{a \in \mathcal{A}} Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t) \quad (4)$$

for the more specific case of SARSA it is:


$$\delta_Q = r_t + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t). \quad (5)$$

# Which action to try?

This is known as the *multi-arm bandit problem* after Las Vegas slot machines

pulling an arm results in a reward for that arm that is random.

Which arm has the best average reward?





## Epsilon-greedy algorithm

Pull the best arm with probability  $1 - \epsilon$   
Pull the other arms, including the best,  
with probability  $\frac{\epsilon}{3}$

A value for  $\epsilon$  might be 10%

# Avoiding obstacles while walking



## State Space

